

# Bubble Sort and its Variants with OpenMP

Sidhartha Ananthula - 21MCME08

November, 2024

## Abstract

This project implements a parallel bubble sort algorithm and an improved parallel bubble sort algorithm using OpenMP. The performance of these parallel algorithms is compared and analyzed.

## 1 System Configuration

The experiments were performed on the following system:

- **Laptop:** MSI GF63
- **Operating System:** openSUSE Leap 15.6
- **CPU:** Intel i5-9300H (8 cores) @ 4.100GHz

## 2 Code Overview

The code is written in C and uses OpenMP, an API for shared-memory MIMD programming. It has three source files: `main.c`, `sort.c`, and `inout.c`. Each file and its functionalities are described below:

### 2.1 `main.c`

This file contains the program's entry point (`main()` function):

- Accepts the number of threads and size of data as command-line arguments.
- Prompts the user to select a sorting algorithm to execute.
- Executes the selected algorithm on randomly generated integer input and displays the execution time.
- Saves sorted data to `output.csv` and execution results to `results.csv`.
- Generated input data is saved in `input.csv`.

### 2.2 `sort.c`

This file contains the implementations of sorting algorithms. Each function takes a pointer to the data and its size as input.

#### 2.2.1 `bubble_sort()`

- Implements the basic Bubble Sort algorithm.

#### 2.2.2 `odd_even_transposition_sort()`

- Implements the Odd-Even Transposition Sort algorithm, as described in [\[Pac22\]](#).

#### 2.2.3 `omp_odd_even_transposition_sort_01()` and `omp_odd_even_transposition_sort_02()`

- Implements the parallel Odd-Even Transposition Sort algorithms using OpenMP, based on the approach in [\[Pac22\]](#).

#### 2.2.4 `shell_sort()`

- Implements the Shell Sort algorithm [She59].

#### 2.2.5 `omp_shell_sort()`

- Implements a parallel version of Shell Sort [Qui88].
- Eliminated loop-carried dependencies by modifying the previous `shell_sort()` algorithm.
- Uses a gradual halving gap sequence [ASA07].
- Each thread sorts sublist(s) of elements, with a reducing gap sequence, using a modified insertion sort.

#### 2.2.6 `omp_shell_sort_02()`

- Implements a parallel version of Shell Sort using the Hibbard Sequence [Qui88].
- The Hibbard Sequence is defined as:

$$g_i = 2^i - 1, \quad i = 1, 2, \dots, \lfloor \log_2(n) \rfloor$$

### 2.3 `inout.c`

#### 2.3.1 `generate_input()`

- Generates random integers in the range  $[1, 0.75 \times \text{size}]$ .
- Saves the data to a CSV file.

#### 2.3.2 `read_input()`

- Reads integer data from a CSV file into a pointer variable.

#### 2.3.3 `write_output()`

- Writes sorted data to a CSV file.

#### 2.3.4 `write_result()`

- Records the algorithm, data size, thread count, and execution time in `results.csv`.
- This data is used for performance analysis.

## 3 Results and Performance Analysis

The initial idea to use a Parallel Shell Sort algorithm to improve performance was inspired from [GGKK03].

The results generated in `results.csv` allow for detailed performance analysis. Multiple runs with varying data sizes and thread counts provide insights into scalability and efficiency.

Each algorithm is run:

- 100 times for the input sizes 1000, 5000, & 10000.
- 50 times for the input sizes 20000, 40000, 60000, & 80000.
- 5 times for the input sizes 100000, & 200000.
- 1 time for the input sizes 300000, 400000, & 500000.
- 5 times for the previous input sizes in the case of parallel shell sort.

Each of the above is run with 1, 2, 4, 8, & 16 threads, in the case of a parallel algorithm.

### 3.1 Serial Algorithms Performance

| Input Size | Execution Time (seconds) |                             |               |
|------------|--------------------------|-----------------------------|---------------|
|            | Bubble Sort              | Odd-Even Transposition Sort | Shell Sort    |
| 1000       | 0.0024028083             | 0.0022502793                | 0.0017570042  |
| 5000       | 0.0616279375             | 0.0558655443                | 0.0490854781  |
| 10000      | 0.2706816557             | 0.2494372873                | 0.1955941217  |
| 20000      | 1.1772869826             | 0.9788865098                | 0.7705120767  |
| 40000      | 4.8420622367             | 3.9837998999                | 3.0417494783  |
| 60000      | 11.0801380679            | 9.1022774343                | 5.6704730329  |
| 80000      | 19.7361797654            | 16.1064127587               | 12.1952081564 |
| 100000     | 30.8882639482            | 25.1810321274               | 16.322841126  |
| 200000     | 124.6022764286           | 100.231252651               | 65.7987660648 |
| 300000     | 280.724216636            | 226.053211357               | 174.246453292 |
| 400000     | 495.188320413            | 400.426751025               | 263.5853297   |
| 500000     | 776.822895897            | 634.396241299               | 394.785282869 |

Table 1: Execution Times for Serial Algorithms

From Table 1 the following can be inferred:

- Shell Sort consistently outperformed Bubble Sort and Odd-Even Transposition Sort for all input sizes, and this is because of its optimized gap-based comparison strategy.
- However, with increase in input size, all 3 serial algorithms show high growth in execution time and hence displays the weakness of serial processing for large data.

### 3.2 Parallel Algorithms Performance

#### Speedup Formula

The speedup  $S(p)$  for an algorithm using  $p$  threads/processes is given by:

$$S(p) = \frac{T(1)}{T(p)}$$

where:

- $T(1)$  is the execution time using 1 thread/process.
  - $T(p)$  is the execution time using  $p$  threads/processes.
- The above Speedup Formula compares the performance improvement when increasing the number of threads from 1 to  $p$ .
  - This is used to calculate the Max Speedup in Tables 2, 3, and 4, to indicate the maximum speedup achieved by that algorithm for a given input size with the use of 2, 4, 8, or 16 threads.

### 3.2.1 Parallel Odd-Even Transposition Sort

| Input Size | Execution Time (seconds) |               |               |               |               | Max Speedup |
|------------|--------------------------|---------------|---------------|---------------|---------------|-------------|
|            | 1 Thread                 | 2 Threads     | 4 Threads     | 8 Threads     | 16 Threads    |             |
| 1000       | 0.0043995729             | 0.0056291712  | 0.0055408402  | 0.0127549767  | 0.0483956481  | 0.79        |
| 5000       | 0.0717712518             | 0.0281106742  | 0.0247649248  | 0.0297316774  | 0.1363136705  | 2.90        |
| 10000      | 0.224461884              | 0.1094497986  | 0.0678424353  | 0.2360771161  | 0.0863311228  | 3.31        |
| 20000      | 0.9341625587             | 0.4431300019  | 0.2716702308  | 0.2743011727  | 0.6438009419  | 3.44        |
| 40000      | 3.9170673179             | 1.8943826517  | 1.0648950346  | 0.9877953404  | 1.7555050802  | 3.97        |
| 60000      | 8.9916148756             | 4.3717881219  | 2.5857709952  | 2.3413656621  | 4.4660238177  | 3.84        |
| 80000      | 16.5034890454            | 8.482658389   | 5.7258598321  | 5.4984473008  | 6.2940302104  | 3.00        |
| 100000     | 25.6711403608            | 12.8538411842 | 7.5483351732  | 7.6591181798  | 9.2750615252  | 3.40        |
| 200000     | 98.8217707366            | 52.5538378264 | 34.4693835264 | 30.6972018662 | 32.5410806188 | 3.22        |
| 300000     | 221.758048915            | 118.154205375 | 75.435695432  | 66.757274992  | 71.407025189  | 3.32        |
| 400000     | 404.49093609             | 210.192399704 | 134.444006517 | 115.949780129 | 123.78182887  | 3.49        |
| 500000     | 643.211717774            | 327.143383715 | 209.064745803 | 185.268861971 | 190.543492433 | 3.47        |

Table 2: Execution Times for Parallel Odd-Even Transposition Sort

From Table 2 the following can be inferred:

- Parallel Odd-Even Transposition sort exhibits a good improvement from the serial version of the same, as the number of threads increases.
- But this still suffers from the high increase in execution time as the input size is increased.
- Another observation that can be made is, for a given input size, the execution time increased from 8 to 16 threads. This is because of the overhead incurred from context-switches among the threads upon trying to use more threads than available cores.
- An improvement with more threads might be visible if executed on a more powerful machine with many more cores available.

### 3.2.2 Parallel Shell Sort - Gradual Halving Gap Sequence

| Input Size | Execution Time (seconds) |              |              |              |              | Max Speedup |
|------------|--------------------------|--------------|--------------|--------------|--------------|-------------|
|            | 1 Thread                 | 2 Threads    | 4 Threads    | 8 Threads    | 16 Threads   |             |
| 1000       | 0.0001160222             | 0.0001258516 | 0.0001378102 | 0.0068138312 | 0.0005858502 | 0.92        |
| 5000       | 0.0007776907             | 0.0005769586 | 0.0005610709 | 0.0098667189 | 0.000895235  | 1.39        |
| 10000      | 0.0015238828             | 0.0009904061 | 0.000743426  | 0.009255718  | 0.001276754  | 2.05        |
| 20000      | 0.003438078              | 0.0020642728 | 0.0014414798 | 0.0071212781 | 0.0025664078 | 2.39        |
| 40000      | 0.007596714              | 0.0043324886 | 0.0030763631 | 0.010429079  | 0.0036892813 | 2.47        |
| 60000      | 0.012006041              | 0.0075229596 | 0.0054520347 | 0.0125844578 | 0.0052426324 | 2.29        |
| 80000      | 0.0172282489             | 0.0099466387 | 0.0071571829 | 0.0198035278 | 0.0069511669 | 2.48        |
| 100000     | 0.024062812              | 0.0128602074 | 0.009147252  | 0.0215467524 | 0.0091316722 | 2.64        |
| 200000     | 0.0546604566             | 0.0291685302 | 0.0194057934 | 0.028510189  | 0.0191293724 | 2.86        |
| 300000     | 0.0823413072             | 0.0437161172 | 0.0328254946 | 0.113351331  | 0.024730842  | 3.33        |
| 400000     | 0.1154088756             | 0.060606625  | 0.0438252642 | 0.107718486  | 0.0440991242 | 2.63        |
| 500000     | 0.1362702092             | 0.0795114286 | 0.0540033234 | 0.2975018922 | 0.046197056  | 2.95        |

Table 3: Execution Times for Parallel Shell Sort (Gradual Halving gap sequence)

From Table 3 the following can be inferred:

- Parallel Shell Sort exhibits a significant improvement in execution time, even for a huge input size of 500,000.
- This algorithm is highly amenable to parallelization and much more efficient to sort huge data than Parallel Odd-Even Transposition sort.
- The overhead when using 16 threads is not observed in Parallel Shell Sort because they are being executed so fast and not much overhead is incurred from context switches.

### 3.2.3 Parallel Shell Sort - Hibbard Gap Sequence

| Input Size | Execution Time (seconds) |              |              |              |              | Max Speedup |
|------------|--------------------------|--------------|--------------|--------------|--------------|-------------|
|            | 1 Thread                 | 2 Threads    | 4 Threads    | 8 Threads    | 16 Threads   |             |
| 1000       | 0.0002227027             | 0.0002270883 | 0.0001796288 | 0.0076201447 | 0.0006867567 | 1.24        |
| 5000       | 0.0008005272             | 0.0005889769 | 0.0004748321 | 0.007167066  | 0.0010458463 | 1.69        |
| 10000      | 0.0018094457             | 0.0011815263 | 0.000940152  | 0.0079271553 | 0.0014104779 | 1.92        |
| 20000      | 0.0041917461             | 0.0025645993 | 0.0018749088 | 0.0138917585 | 0.0023091777 | 2.24        |
| 40000      | 0.0093606897             | 0.0055948565 | 0.003891705  | 0.011296769  | 0.0036881603 | 2.54        |
| 60000      | 0.0147576036             | 0.0081541226 | 0.0054454264 | 0.0094921519 | 0.0053113553 | 2.78        |
| 80000      | 0.0197075926             | 0.0118485022 | 0.0080760815 | 0.009757968  | 0.007586132  | 2.60        |
| 100000     | 0.0284322196             | 0.014945354  | 0.0099236306 | 0.0078345594 | 0.0089508908 | 3.63        |
| 200000     | 0.0586942048             | 0.0326963984 | 0.020659978  | 0.0150368526 | 0.0172142492 | 3.90        |
| 300000     | 0.0954365884             | 0.0501971524 | 0.0363003164 | 0.1515741846 | 0.029540666  | 3.23        |
| 400000     | 0.1323158122             | 0.0746978842 | 0.048753065  | 0.1581163424 | 0.040731122  | 3.25        |
| 500000     | 0.1731763882             | 0.0999991686 | 0.0663707694 | 0.1987161396 | 0.050199827  | 3.45        |

Table 4: Execution Times for Parallel Shell Sort (Hibbard gap sequence)

From Table 4 the following can be inferred:

- This version of Parallel Shell sort provided similar results to the previous version with a significant improvement in execution times.
- The previous version performs, not much but, slightly better than this in many cases, as it was suggested an Enhanced Shell Sort [ASA07].

## 3.3 Graphical Analysis

### 3.3.1 Plot between Parallel Shell Sort and Parallel Odd-Even Transposition Sort

Here we compare the performance of Parallel Shell Sort and Parallel Odd-Even Transposition Sort. The plot shows the time complexity of both algorithms under various input sizes using 8 threads. It turns out that Parallel Shell Sort is highly efficient.




Figure 1: Comparison of Parallel Shell Sort and Parallel Odd-Even Transposition Sort

Speedup plot for the above:




Figure 2: Speedup of Parallel Odd-Even Transposition Sort vs Parallel Shell Sort

### 3.3.2 Plot between Odd-Even Transposition Sort and Parallel Odd-Even Transposition Sort

This plot shows the comparison between the classic Odd-Even Transposition Sort and its parallel variant run with 8 threads. But neither are scalable.




Figure 3: Comparison of Odd-Even Transposition Sort and Parallel Odd-Even Transposition Sort

Speedup plot for the above:




Figure 4: Speedup of Odd-Even Transposition Sort vs Parallel Odd-Even Transposition Sort

### 3.3.3 Plot between Shell Sort and Parallel Shell Sort variants

This figure compares the performance of the sequential Shell Sort algorithm and its parallel implementations with 8 threads. It highlights how parallelization enhances the algorithm's high efficiency and scalability with larger datasets.




Figure 5: Comparison of Shell Sort and Parallel Shell Sort

Speedup plot for the above:




Figure 6: Speedup of Shell Sort vs Parallel Shell Sort variants

## References

- [ASA07] Tanveer Afzal, Basit Shahzad, and Salman Aslam. Enhanced shell sort algorithm. In *International Journal of Computer, Electrical, Automation, Control and Information Engineering Vol:1, No:3, 2007*, 02 2007.
- [GGKK03] Ananth Grama, Anshul Gupta, George Karypis, and Vipin Kumar. *Introduction to Parallel Computing*. Addison Wesley, 2003.
- [Pac22] Peter S. Pacheco. *An Introduction to Parallel Programming*. Morgan Kaufmann, 2022.
- [Qui88] Michael J. Quinn. Parallel sorting algorithms for tightly coupled multiprocessors. *Parallel Computing*, 6(3):349–357, 1988.
- [She59] D. L. Shell. A high-speed sorting procedure. *Commun. ACM*, 2(7):30–32, July 1959.